

```
In [2]: import numpy as np  
import pandas as pd
```

```
In [3]: housing= pd.read_csv("train (1).csv")
```

```
In [1]: #print(housing.to_string())
```

In [5]: `housing.info()`

```
<class 'pandas.core.frame.DataFrame'>
RangeIndex: 1460 entries, 0 to 1459
Data columns (total 81 columns):
#   Column                Non-Null Count  Dtype
---  -
0   Id                    1460 non-null   int64
1   MSSubClass            1460 non-null   int64
2   MSZoning              1460 non-null   object
3   LotFrontage          1201 non-null   float64
4   LotArea               1460 non-null   int64
5   Street               1460 non-null   object
6   Alley                91 non-null     object
7   LotShape              1460 non-null   object
8   LandContour          1460 non-null   object
9   Utilities            1460 non-null   object
10  LotConfig             1460 non-null   object
11  LandSlope             1460 non-null   object
12  Neighborhood          1460 non-null   object
13  Condition1            1460 non-null   object
14  Condition2            1460 non-null   object
15  BldgType              1460 non-null   object
16  HouseStyle            1460 non-null   object
17  OverallQual           1460 non-null   int64
18  OverallCond           1460 non-null   int64
19  YearBuilt             1460 non-null   int64
20  YearRemodAdd          1460 non-null   int64
21  RoofStyle             1460 non-null   object
22  RoofMatl              1460 non-null   object
23  Exterior1st           1460 non-null   object
24  Exterior2nd           1460 non-null   object
25  MasVnrType            588 non-null     object
26  MasVnrArea            1452 non-null   float64
27  ExterQual             1460 non-null   object
28  ExterCond             1460 non-null   object
29  Foundation            1460 non-null   object
30  BsmtQual              1423 non-null   object
31  BsmtCond              1423 non-null   object
32  BsmtExposure          1422 non-null   object
33  BsmtFinType1          1423 non-null   object
34  BsmtFinSF1            1460 non-null   int64
35  BsmtFinType2          1422 non-null   object
36  BsmtFinSF2            1460 non-null   int64
37  BsmtUnfSF             1460 non-null   int64
38  TotalBsmtSF           1460 non-null   int64
39  Heating               1460 non-null   object
40  HeatingQC             1460 non-null   object
41  CentralAir            1460 non-null   object
42  Electrical            1459 non-null   object
43  1stFlrSF              1460 non-null   int64
44  2ndFlrSF              1460 non-null   int64
45  LowQualFinSF          1460 non-null   int64
46  GrLivArea             1460 non-null   int64
47  BsmtFullBath          1460 non-null   int64
48  BsmtHalfBath          1460 non-null   int64
49  FullBath              1460 non-null   int64
50  HalfBath              1460 non-null   int64
51  BedroomAbvGr          1460 non-null   int64
52  KitchenAbvGr          1460 non-null   int64
53  KitchenQual           1460 non-null   object
54  TotRmsAbvGrd          1460 non-null   int64
55  Functional            1460 non-null   object
```

```

56 Fireplaces      1460 non-null    int64
57 FireplaceQu     770 non-null    object
58 GarageType      1379 non-null   object
59 GarageYrBltd    1379 non-null   float64
60 GarageFinish    1379 non-null   object
61 GarageCars      1460 non-null   int64
62 GarageArea      1460 non-null   int64
63 GarageQual      1379 non-null   object
64 GarageCond      1379 non-null   object
65 PavedDrive      1460 non-null   object
66 WoodDeckSF      1460 non-null   int64
67 OpenPorchSF     1460 non-null   int64
68 EnclosedPorch   1460 non-null   int64
69 3SsnPorch       1460 non-null   int64
70 ScreenPorch     1460 non-null   int64
71 PoolArea        1460 non-null   int64
72 PoolQC          7 non-null      object
73 Fence           281 non-null    object
74 MiscFeature     54 non-null     object
75 MiscVal         1460 non-null   int64
76 MoSold          1460 non-null   int64
77 YrSold           1460 non-null   int64
78 SaleType        1460 non-null   object
79 SaleCondition   1460 non-null   object
80 SalePrice       1460 non-null   int64
dtypes: float64(3), int64(35), object(43)
memory usage: 924.0+ KB


```

In [6]: `housing.describe()`

Out[6]:

	Id	MSSubClass	LotFrontage	LotArea	OverallQual	OverallCond	
count	1460.000000	1460.000000	1201.000000	1460.000000	1460.000000	1460.000000	1460
mean	730.500000	56.897260	70.049958	10516.828082	6.099315	5.575342	1973
std	421.610009	42.300571	24.284752	9981.264932	1.382997	1.112799	3.562204
min	1.000000	20.000000	21.000000	1300.000000	1.000000	1.000000	1871
25%	365.750000	20.000000	59.000000	7553.500000	5.000000	5.000000	1961
50%	730.500000	50.000000	69.000000	9478.500000	6.000000	5.000000	1973
75%	1095.250000	70.000000	80.000000	11601.500000	7.000000	6.000000	2063
max	1460.000000	190.000000	313.000000	215245.000000	10.000000	9.000000	2063

8 rows × 38 columns



```
In [10]: # Calculate the percentage of missing values in each column
missing_percentage = housing.isna().mean() * 100

# Filter columns with missing values
missing_columns = missing_percentage[missing_percentage > 0]

# Output the filtered columns and the count of them
missing_columns, len(missing_columns)
```

```
Out[10]: (LotFrontage      17.739726
Alley                    93.767123
MasVnrType              59.726027
MasVnrArea              0.547945
BsmtQual                2.534247
BsmtCond                2.534247
BsmtExposure            2.602740
BsmtFinType1            2.534247
BsmtFinType2            2.602740
Electrical              0.068493
FireplaceQu            47.260274
GarageType              5.547945
GarageYrBlt             5.547945
GarageFinish            5.547945
GarageQual              5.547945
GarageCond              5.547945
PoolQC                 99.520548
Fence                   80.753425
MiscFeature            96.301370
dtype: float64,
19)
```


```
In [11]: # Forward fill missing values for object-type columns with missing data
housing.loc[:, missing_columns.reset_index()['index']].select_dtypes(include=object).fillna(method='ffill')

# Backward fill missing values for object-type columns with missing data
housing.loc[:, missing_columns.reset_index()['index']].select_dtypes(include=object).fillna(method='bfill')
```

```
Out[11]:
```

	Alley	MasVnrType	BsmtQual	BsmtCond	BsmtExposure	BsmtFinType1	BsmtFinType2
0	Grvl	BrkFace	Gd	TA	No	GLQ	Unf
1	Grvl	BrkFace	Gd	TA	Gd	ALQ	Unf
2	Grvl	BrkFace	Gd	TA	Mn	GLQ	Unf
3	Grvl	BrkFace	TA	Gd	No	ALQ	Unf
4	Grvl	BrkFace	Gd	TA	Av	GLQ	Unf
...
1455	NaN	Stone	Gd	TA	No	Unf	Unf
1456	NaN	Stone	Gd	TA	No	ALQ	Refr
1457	NaN	NaN	TA	Gd	No	GLQ	Unf
1458	NaN	NaN	TA	TA	Mn	GLQ	Refr
1459	NaN	NaN	TA	TA	No	BLQ	LwC

1460 rows × 16 columns



```
In [12]: # Count the missing values in each column
missing_values = housing.isna().sum()

# Output the missing values count for each column
print(missing_values)
```

```
Id                0
MSSubClass        0
MSZoning          0
LotFrontage      259
LotArea          0
...
MoSold           0
YrSold           0
SaleType         0
SaleCondition    0
SalePrice        0
Length: 81, dtype: int64
```

```
In [13]: # Calculate the mean of float64-type columns with missing data
housing.loc[:, missing_columns.reset_index()['index']].select_dtypes(include=float).mean()
```

```
Out[13]: LotFrontage      70.049958
MasVnrArea      103.685262
GarageYrBlt     1978.506164
dtype: float64
```

```
In [14]: # Impute missing values in 'LotFrontage' with a default value of 70.04
housing['LotFrontage'] = np.where(housing['LotFrontage'].isna(), 70.04, housing['LotFrontage'])

# Impute missing values in 'MasVnrArea' with a default value of 103.685262
housing['MasVnrArea'] = np.where(housing['MasVnrArea'].isna(), 103.685262, housing['MasVnrArea'])


# Impute missing values in 'GarageYrBlt' with a default value of 1978.50616
housing['GarageYrBlt'] = np.where(housing['GarageYrBlt'].isna(), 1978.50616, housing['GarageYrBlt'])
```

```
In [15]: housing
```

```
Out[15]:
```

	Id	MSSubClass	MSZoning	LotFrontage	LotArea	Street	Alley	LotShape	LandCont
0	1	60	RL	65.0	8450	Pave	NaN	Reg	
1	2	20	RL	80.0	9600	Pave	NaN	Reg	
2	3	60	RL	68.0	11250	Pave	NaN	IR1	
3	4	70	RL	60.0	9550	Pave	NaN	IR1	
4	5	60	RL	84.0	14260	Pave	NaN	IR1	
...	
1455	1456	60	RL	62.0	7917	Pave	NaN	Reg	
1456	1457	20	RL	85.0	13175	Pave	NaN	Reg	
1457	1458	70	RL	66.0	9042	Pave	NaN	Reg	
1458	1459	20	RL	68.0	9717	Pave	NaN	Reg	
1459	1460	20	RL	75.0	9937	Pave	NaN	Reg	

1460 rows × 81 columns



```
In [16]: # Count the missing values in each column
missing_values = housing.isna().sum()

# Drop columns with more than 50% missing values
housing.drop(missing_values[missing_values > 50].reset_index()['index'], axis=1)

# Output the shape of the dataframe after column removal
housing.shape
```

```
Out[16]: (1460, 71)
```

Converting cat cols into numerical cols

```
In [17]: # Get the column names with 'object' data type (categorical columns)
df_cat=housing.select_dtypes(include=['object']).columns
```

In [18]: df_cat

```
Out[18]: Index(['MSZoning', 'Street', 'LotShape', 'LandContour', 'Utilities',
               'LotConfig', 'LandSlope', 'Neighborhood', 'Condition1', 'Condition
               2',
               'BldgType', 'HouseStyle', 'RoofStyle', 'RoofMatl', 'Exterior1st',
               'Exterior2nd', 'ExterQual', 'ExterCond', 'Foundation', 'BsmtQual',
               'BsmtCond', 'BsmtExposure', 'BsmtFinType1', 'BsmtFinType2', 'Heatin
               g',
               'HeatingQC', 'CentralAir', 'Electrical', 'KitchenQual', 'Functiona
               l',
               'PavedDrive', 'SaleType', 'SaleCondition'],
              dtype='object')
```

```
In [21]: # Importing necessary libraries for preprocessing and metrics
from sklearn import preprocessing, metrics
from sklearn.preprocessing import OneHotEncoder
from sklearn.compose import ColumnTransformer
from sklearn.impute import SimpleImputer
```

```
In [22]: # Initialize the OneHotEncoder for categorical data encoding
ohe = OneHotEncoder()
```

```
In [23]: # Reset the index of the DataFrame without keeping the old index
housing.reset_index(drop=True, inplace=True)

# List of categorical columns in the housing dataset
df_cat_cols = ['MSZoning', 'Street', 'Alley', 'LotShape', 'LandContour', 'U
               'LotConfig', 'LandSlope', 'Neighborhood', 'Condition1', 'Con
               'BldgType', 'HouseStyle', 'RoofStyle', 'RoofMatl', 'Exterior
               'Exterior2nd', 'MasVnrType', 'ExterQual', 'ExterCond', 'Foun
               'BsmtQual', 'BsmtCond', 'BsmtExposure', 'BsmtFinType1', 'Bsm
               'Heating', 'HeatingQC', 'CentralAir', 'Electrical', 'Kitchen
               'Functional', 'FireplaceQu', 'GarageType', 'GarageFinish', '
               'GarageCond', 'PavedDrive', 'PoolQC', 'Fence', 'MiscFeature'
               'SaleCondition']
```

```
In [24]: # Apply OneHotEncoder to the categorical columns in housing dataset
ohe_var = ohe.fit_transform(housing[['MSZoning', 'Street', 'LotShape', 'Lan
               'LotConfig', 'LandSlope', 'Neighborhood
               'BldgType', 'HouseStyle', 'RoofStyle',
               'Exterior2nd', 'ExterQual', 'ExterCond
               'BsmtCond', 'BsmtExposure', 'BsmtFinTy
               'HeatingQC', 'CentralAir', 'Electrical
               'PavedDrive', 'SaleType', 'SaleCondi
```



```
In [25]: # Convert the OneHotEncoded array into a DataFrame with appropriate column  
ohe_var = pd.DataFrame(ohe_var.toarray(), columns=ohe.get_feature_names_out)
```

```
In [26]: # Concatenate the original housing DataFrame with the OneHotEncoded columns  
housing_new = pd.concat([housing, ohe_var], axis=1)
```

```
In [2]: #housing_new
```

```
In [28]: # Drop the original categorical columns after OneHotEncoding  
housing_new.drop(columns=['MSZoning', 'Street', 'LotShape', 'LandContour',  
                           'LotConfig', 'LandSlope', 'Neighborhood', 'Condit  
                           'BldgType', 'HouseStyle', 'RoofStyle', 'RoofMatl'  
                           'Exterior2nd', 'ExterQual', 'ExterCond', 'Foundat  
                           'BsmtCond', 'BsmtExposure', 'BsmtFinType1', 'Bsmt  
                           'HeatingQC', 'CentralAir', 'Electrical', 'Kitchen  
                           'PavedDrive', 'SaleType', 'SaleCondition'], inpla
```

```
In [3]: #housing_new
```

Type *Markdown* and LaTeX: α^2

```
In [36]: from statsmodels.stats.outliers_influence import variance_inflation_factor
from statsmodels.tools.tools import add_constant

# Add a constant term to the DataFrame to account for the intercept in VIF
df_constant = add_constant(housing_new)

# Calculate VIF for each variable
vif_data = pd.DataFrame()
vif_data["Feature"] = df_constant.columns
vif_data["VIF"] = [variance_inflation_factor(df_constant.values, i) for i in range(df_constant.shape[0])]

# Display the VIF values for each feature
print(vif_data)
```

C:\ProgramData\anaconda3\Lib\site-packages\statsmodels\regression\linear_model.py:1781: RuntimeWarning: divide by zero encountered in scalar divide

return 1 - self.ssr/self.centered_tss

C:\ProgramData\anaconda3\Lib\site-packages\statsmodels\stats\outliers_influence.py:198: RuntimeWarning: divide by zero encountered in scalar divide

vif = 1. / (1. - r_squared_i)

	Feature	VIF
0	const	0.000000
1	Id	1.192539
2	MSSubClass	33.805402
3	LotFrontage	2.505810
4	LotArea	3.417143
..
251	SaleCondition_AdjLand	inf
252	SaleCondition_Alloca	inf
253	SaleCondition_Family	inf
254	SaleCondition_Normal	inf

```
In [37]: # Select columns with VIF values greater than or equal to 20
reqcols = vif_data.loc[vif_data['VIF'] >= 20, 'Feature']

# Access the corresponding columns in the housing_new DataFrame
housing_new.loc[:, reqcols]
```

Out[37]:

	MSSubClass	BsmtFinSF1	BsmtFinSF2	BsmtUnfSF	TotalBsmtSF	1stFlrSF	2ndFlrSF
0	60	706	0	150	856	856	854
1	20	978	0	284	1262	1262	0
2	60	486	0	434	920	920	866
3	70	216	0	540	756	961	756
4	60	655	0	490	1145	1145	1053
...
1455	60	0	0	953	953	953	694
1456	20	790	163	589	1542	2073	0
1457	70	275	0	877	1152	1188	1152
1458	20	49	1029	0	1078	1078	0
1459	20	830	290	136	1256	1256	0

1460 rows × 226 columns



```
In [38]: def correlation(df, threshold):
    correlated_cols = set()
    corre_matrix = df.corr()
    for i in range(len(corre_matrix.columns)):
        for j in range(i):
            if abs(corre_matrix.iloc[i, j]) > threshold:
                correlated_cols.add(corre_matrix.columns[i])
    return correlated_cols
```

```
In [39]: # Function to remove highly correlated features (multicollinearity)
def remove_multicollinearity(df, threshold):
    correlated_cols = correlation(df, threshold)

    # Drop highly correlated columns
    df_cleaned = df.drop(columns=correlated_cols)

    return df_cleaned

# Apply the function with a threshold (e.g., 0.9 for high correlation)
threshold = 0.9
housing_cleaned = remove_multicollinearity(housing_new, threshold)

# Check the shape of the new DataFrame after removing correlated features
print(housing_cleaned.shape)
```

(1460, 239)

```
In [40]: # Remove multicollinearity using correlation before training the model
threshold = 0.9
housing_cleaned = remove_multicollinearity(housing_new, threshold)

# Now, split the data into training and test sets
X = housing_cleaned.loc[:, housing_cleaned.columns != 'SalePrice'].values
y = housing_cleaned.loc[:, housing_cleaned.columns == 'SalePrice'].values

# Train/Test Split
X_train, X_test, y_train, y_test = train_test_split(X, y, test_size=0.2, ra

# Model Training
regressor = LinearRegression()
regressor.fit(X_train, y_train)
```

Out[40]: LinearRegression()

In a Jupyter environment, please rerun this cell to show the HTML representation or trust the notebook.

On GitHub, the HTML representation is unable to render, please try loading this page with nbviewer.org.

```
In [41]: from sklearn.metrics import mean_absolute_error, mean_squared_error, r2_score

# Predict on the test set
y_pred = regressor.predict(X_test)

# Evaluate performance
mae = mean_absolute_error(y_test, y_pred)
mse = mean_squared_error(y_test, y_pred)
rmse = mean_squared_error(y_test, y_pred, squared=False)
r2 = r2_score(y_test, y_pred)

# Print evaluation metrics
print("Mean Absolute Error:", mae)
print("Mean Squared Error:", mse)
print("Root Mean Squared Error:", rmse)
print("R² Score:", r2)
```

Mean Absolute Error: 22259.56716756969
Mean Squared Error: 3096707469.952557
Root Mean Squared Error: 55648.067980412015
R² Score: 0.5515819940928334

```
In [42]: # Predict on the training set
y_train_pred = regressor.predict(X_train)

# Compare training vs test performance
train_mae = mean_absolute_error(y_train, y_train_pred)
test_mae = mean_absolute_error(y_test, y_pred)

train_r2 = r2_score(y_train, y_train_pred)
test_r2 = r2_score(y_test, y_pred)

# Print comparison
print("Training MAE:", train_mae)
print("Test MAE:", test_mae)
print("Training R²:", train_r2)
print("Test R²:", test_r2)
```

Training MAE: 12548.10328137777
Test MAE: 22259.56716756969
Training R²: 0.94058573886167
Test R²: 0.5515819940928334

```
In [43]: from sklearn.model_selection import GridSearchCV
from sklearn.linear_model import Ridge

# Hyperparameter tuning for Ridge regression (example)
ridge = Ridge()
param_grid = {'alpha': [0.1, 1, 10, 100]} # Tuning the regularization para
grid_search = GridSearchCV(ridge, param_grid, cv=5)
grid_search.fit(X_train, y_train)

# Best hyperparameters
print("Best alpha:", grid_search.best_params_)

# Train and evaluate with the best model
best_ridge = grid_search.best_estimator_
y_pred_best = best_ridge.predict(X_test)
print("Best Model R²:", r2_score(y_test, y_pred_best))
```

Best alpha: {'alpha': 10}
 Best Model R²: 0.7093565606475368

```
In [44]: #adjusted R2
# Function to compute adjusted R²
def adjusted_r2_score(X, y, model):
    n = len(y) # Number of observations
    p = X.shape[1] # Number of features
    r2 = r2_score(y, model.predict(X)) # R² Score
    adj_r2 = 1 - (1 - r2) * (n - 1) / (n - p - 1)
    return adj_r2

# Calculate Adjusted R² for both models
train_adj_r2 = adjusted_r2_score(X_train, y_train, regressor)
test_adj_r2 = adjusted_r2_score(X_test, y_test, regressor)

print("Training Adjusted R²:", train_adj_r2)
print("Test Adjusted R²:", test_adj_r2)
```

Training Adjusted R²: 0.9253644319177274
 Test Adjusted R²: -1.4620686739431226

```
In [45]: from sklearn.linear_model import Ridge

ridge = Ridge(alpha=10) # You can start with alpha=10, but we can tune it
ridge.fit(X_train, y_train)
y_pred_ridge = ridge.predict(X_test)
print("R² (Ridge):", r2_score(y_test, y_pred_ridge))
```

R² (Ridge): 0.7093565606475368

```
In [46]: from sklearn.preprocessing import StandardScaler

# Scaling the features
scaler = StandardScaler()
X_train_scaled = scaler.fit_transform(X_train)
X_test_scaled = scaler.transform(X_test)

# Fit Ridge on scaled data
ridge.fit(X_train_scaled, y_train)
y_pred_ridge_scaled = ridge.predict(X_test_scaled)
print("R² (Ridge with Scaling):", r2_score(y_test, y_pred_ridge_scaled))
```

R² (Ridge with Scaling): 0.5878693901662428

```
In [47]: from sklearn.model_selection import cross_val_score

# Use Ridge with cross-validation
cv_scores = cross_val_score(ridge, X_train_scaled, y_train, cv=5, scoring='neg_mean_squared_error')
print(f"Cross-Validation MSE (negative): {cv_scores}")
```

Cross-Validation MSE (negative): [-7.23414763e+08 -1.48384541e+09 -5.58768417e+08 -8.23647330e+08 -4.14967786e+08]

```
In [48]: from sklearn.model_selection import GridSearchCV

param_grid = {'alpha': [0.1, 1, 10, 100, 1000]}
grid_search = GridSearchCV(ridge, param_grid, cv=5, scoring='neg_mean_squared_error')
grid_search.fit(X_train_scaled, y_train)

print(f"Best parameters: {grid_search.best_params_}")
```

Best parameters: {'alpha': 100}

```
In [49]: # Train with the best alpha found from GridSearchCV
best_ridge = Ridge(alpha=100)
best_ridge.fit(X_train_scaled, y_train)
y_pred_best_ridge = best_ridge.predict(X_test_scaled)

print("R² (Ridge with alpha=100):", r2_score(y_test, y_pred_best_ridge))
```

R² (Ridge with alpha=100): 0.6307215979489333

```
In [50]: from sklearn.linear_model import Lasso

lasso = Lasso(alpha=100)
lasso.fit(X_train_scaled, y_train)
y_pred_lasso = lasso.predict(X_test_scaled)

print("R² (Lasso with alpha=100):", r2_score(y_test, y_pred_lasso))
```

R² (Lasso with alpha=100): 0.5968935100449475

```
In [58]: from sklearn.metrics import mean_squared_error, mean_absolute_error, r2_score

# Make predictions using the trained Random Forest model
y_pred = rf_model.predict(X_train)

# Evaluate performance
mse = mean_squared_error(y_train, y_pred)
mae = mean_absolute_error(y_train, y_pred)
r2 = r2_score(y_train, y_pred)

# Print evaluation metrics
print(f"Mean Squared Error (MSE): {mse}")
print(f"Mean Absolute Error (MAE): {mae}")
print(f"R2 (Coefficient of Determination): {r2}")
```

Mean Squared Error (MSE): 127450989.59156996
Mean Absolute Error (MAE): 6522.615590753425
R² (Coefficient of Determination): 0.97929968918002

```
In [59]: # Predicting on the test set
y_test_pred = rf_model.predict(X_test)

# Evaluating the model performance on the test set
from sklearn.metrics import mean_squared_error, mean_absolute_error, r2_score

# Mean Squared Error
mse_test = mean_squared_error(y_test, y_test_pred)
# Mean Absolute Error
mae_test = mean_absolute_error(y_test, y_test_pred)
# R2
r2_test = r2_score(y_test, y_test_pred)

print("Test Mean Squared Error (MSE):", mse_test)
print("Test Mean Absolute Error (MAE):", mae_test)
print("Test R2:", r2_test)
```

Test Mean Squared Error (MSE): 1196368366.8795376
Test Mean Absolute Error (MAE): 17709.577739726028
Test R²: 0.8267601565172202


```
In [60]: def adjusted_r2(R2, n, p):  
         return 1 - (1 - R2) * (n - 1) / (n - p - 1)  
  
         # Calculate Adjusted R2 for the training set  
         n_train = X_train.shape[0] # Number of samples in the training set  
         p_train = X_train.shape[1] # Number of features in the training set  
         adj_r2_train = adjusted_r2(rf_model.score(X_train, y_train), n_train, p_train)  
  
         # Calculate Adjusted R2 for the test set  
         n_test = X_test.shape[0] # Number of samples in the test set  
         p_test = X_test.shape[1] # Number of features in the test set  
         adj_r2_test = adjusted_r2(rf_model.score(X_test, y_test), n_test, p_test)  
  
         # Display the Adjusted R2 values  
         print("Adjusted R2 (Train):", adj_r2_train)  
         print("Adjusted R2 (Test):", adj_r2_test)
```

Adjusted R² (Train): 0.973996487915052
Adjusted R² (Test): 0.048815198990775244

```
In [61]: from sklearn.preprocessing import StandardScaler  
  
         scaler = StandardScaler()  
         X_scaled = scaler.fit_transform(X_train) # Scale the train data  
         X_test_scaled = scaler.transform(X_test) # Use the same scaling for test data
```

```
In [62]: from sklearn.decomposition import PCA  
  
         # Initialize PCA and fit it to the scaled data  
         pca = PCA(n_components=0.95) # Keep 95% of variance  
         X_train_pca = pca.fit_transform(X_scaled)  
         X_test_pca = pca.transform(X_test_scaled)  
  
         print(f"Original number of features: {X_scaled.shape[1]}")  
         print(f"Reduced number of features after PCA: {X_train_pca.shape[1]}")
```

Original number of features: 238
Reduced number of features after PCA: 153

```
In [63]: from sklearn.ensemble import RandomForestRegressor

# Initialize and train the Random Forest model
rf_pca = RandomForestRegressor(random_state=42)
rf_pca.fit(X_train_pca, y_train)

# Evaluate the model
train_pred_pca = rf_pca.predict(X_train_pca)
test_pred_pca = rf_pca.predict(X_test_pca)

# Calculate the performance metrics
from sklearn.metrics import mean_squared_error, mean_absolute_error, r2_score

# Train metrics
train_mse = mean_squared_error(y_train, train_pred_pca)
train_mae = mean_absolute_error(y_train, train_pred_pca)
train_r2 = r2_score(y_train, train_pred_pca)

# Test metrics
test_mse = mean_squared_error(y_test, test_pred_pca)
test_mae = mean_absolute_error(y_test, test_pred_pca)
test_r2 = r2_score(y_test, test_pred_pca)

# Print results
print(f"Train MSE: {train_mse}")
print(f"Train MAE: {train_mae}")
print(f"Train R²: {train_r2}")

print(f"Test MSE: {test_mse}")
print(f"Test MAE: {test_mae}")
print(f"Test R²: {test_r2}")
```

C:\ProgramData\anaconda3\Lib\site-packages\sklearn\base.py:1151: DataConversionWarning: A column-vector y was passed when a 1d array was expected. Please change the shape of y to (n_samples,), for example using ravel().

```
    return fit_method(estimator, *args, **kwargs)
```

```
Train MSE: 135271553.29757774
Train MAE: 7101.626917808219
Train R²: 0.9780294903371503
Test MSE: 1520675043.5950634
Test MAE: 20648.582226027396
Test R²: 0.7797990035228807
```

```
In [64]: from sklearn.model_selection import GridSearchCV
from sklearn.ensemble import RandomForestRegressor

# Define the Random Forest model
rf = RandomForestRegressor(random_state=42)

# Set up the parameter grid to search through
param_grid = {
    'n_estimators': [50, 100, 200],
    'max_depth': [None, 10, 20, 30],
    'min_samples_split': [2, 5, 10],
    'min_samples_leaf': [1, 2, 4],
    'bootstrap': [True, False]
}

# Set up the grid search with cross-validation
grid_search = GridSearchCV(estimator=rf, param_grid=param_grid,
                           cv=5, n_jobs=-1, verbose=2, scoring='neg_mean_sq

# Fit the grid search to the training data
grid_search.fit(X_train, y_train)

# Get the best parameters from the grid search
best_params = grid_search.best_params_

# Get the best model from the grid search
best_rf_model = grid_search.best_estimator_

# Evaluate the best model on the test data
y_test_pred = best_rf_model.predict(X_test)

# Print best parameters and evaluation metrics
print("Best parameters:", best_params)

from sklearn.metrics import mean_squared_error, mean_absolute_error, r2_score

test_mse = mean_squared_error(y_test, y_test_pred)
test_mae = mean_absolute_error(y_test, y_test_pred)
test_r2 = r2_score(y_test, y_test_pred)

print(f'Test MSE: {test_mse}')
print(f'Test MAE: {test_mae}')
print(f'Test R²: {test_r2}')
```

Fitting 5 folds for each of 216 candidates, totalling 1080 fits

C:\ProgramData\anaconda3\Lib\site-packages\sklearn\base.py:1151: DataConversionWarning: A column-vector y was passed when a 1d array was expected. Please change the shape of y to (n_samples,), for example using ravel().

```
return fit_method(estimator, *args, **kwargs)
```

```
Best parameters: {'bootstrap': True, 'max_depth': 20, 'min_samples_leaf': 1, 'min_samples_split': 2, 'n_estimators': 200}
Test MSE: 111590026.6556997
Test MAE: 17525.504969436388
Test R²: 0.8384123558327681
```

```
In [4]: import matplotlib.pyplot as plt


# Define models and their R2 scores (replace these values with actual results)
models = ['Ridge', 'Lasso', 'Random Forest', 'PCA + Random Forest', 'Random Forest (Tuned)']
train_r2_scores = [0.7094, 0.5969, 0.9793, 0.9780, 0.9740] # Add Train R2
test_r2_scores = [0.6307, 0.5169, 0.8268, 0.7798, 0.8384] # Add Test R2

# Create bar width and positions
x_pos = range(len(models))
width = 0.4

# Plot
plt.figure(figsize=(10, 6))
plt.bar(x=[p - width/2 for p in x_pos], height=train_r2_scores, width=width, color='lightblue')
plt.bar(x=[p + width/2 for p in x_pos], height=test_r2_scores, width=width, color='salmon')

# Add labels and title
plt.xlabel('Models', fontsize=12)
plt.ylabel('R2 Score', fontsize=12)
plt.title('Model Accuracy Comparison (R2 Score)', fontsize=14)
plt.xticks(x_pos, models, rotation=15)
plt.legend()
plt.tight_layout()

# Show Plot
plt.show()
```



```
In [ ]:
```